



A black box for online approximate pattern matching[☆]

Raphaël Clifford^{a,*}, Klim Efremenko^{b,c}, Benny Porat^b, Ely Porat^b

^a University of Bristol, Dept. of Computer Science, Bristol BS8 1UB, UK

^b Bar-Ilan University, Dept. of Computer Science, 52900 Ramat-Gan, Israel

^c Weizman Institute, Dept. of Computer Science and Applied Mathematics, Rehovot, Israel

ARTICLE INFO

Article history:

Received 9 April 2008

Revised 21 May 2010

Available online 25 January 2011

ABSTRACT

We present a deterministic black box solution for online approximate matching. Given a pattern of length m and a streaming text of length n that arrives one character at a time, the task is to report the distance between the pattern and a sliding window of the text as soon as the new character arrives. Our solution requires $O(\sum_{j=1}^{\log_2 m} T(n, 2^{j-1})/n)$ time for each input character, where $T(n, m)$ is the total running time of the best offline algorithm. The types of approximation that are supported include exact matching with wildcards, matching under the Hamming norm, approximating the Hamming norm, k -mismatch and numerical measures such as the L_2 and L_1 norms. For these examples, the resulting online algorithms take $O(\log^2 m)$, $O(\sqrt{m \log m})$, $O(\log^2 m / \epsilon^2)$, $O(\sqrt{k \log k \log m})$, $O(\log^2 m)$ and $O(\sqrt{m \log m})$ time per character, respectively. The space overhead is linear in the pattern size, which we show is optimal for any deterministic algorithm.

© 2011 Elsevier Inc. All rights reserved.

1. Introduction

Fast approximate string matching is a central problem of modern data intensive applications. Its applications are many and varied, from computational biology and large scale web searching to searching multimedia databases and digital libraries. As a result, string matching has to continuously adapt itself to the problem at hand. Simultaneously, the need for asymptotically fast algorithms grows every year with the explosion of data available in digital form.

A great deal of progress has been made in finding fast algorithms for a variety of important forms of approximate matching. One of the most studied is the Hamming distance which measures the number of mismatches between two strings. Given a text t of length n and a pattern p of length m , the task is to report the Hamming distance at every possible alignment. $O(n\sqrt{m \log m})$ time solutions based on repeated applications of the FFT were given independently by both Abrahamson and Kosaraju in 1987 [1,19]. Particular interest has been paid to a bounded version of this problem called the k -mismatch problem. Here a bound k is given and we need only report the Hamming distance if it is less than or equal to k . If the number of mismatches is greater than the bound, the algorithm need only report that fact and not give the actual Hamming distance. In 1985, Landau and Vishkin gave a beautiful $O(nk)$ algorithm that is not FFT based and uses constant time LCA (Lowest Common Ancestor) operations on the suffix tree of p and t [20]. This was subsequently improved to $O(n\sqrt{k \log k})$ time by a method based on filtering and FFTs again [3]. A separate line of research considered the question of how to find approximations within a $(1 + \epsilon)$ multiplicative factor of the Hamming distance [16,18].

The problem of determining the time complexity of exact matching with do not cares has also been well studied over many years [6,11,13,16,17], culminating in two related deterministic $O(n \log m)$ time solutions. This has been accompanied

[☆] This work was partly funded by the BSF, ISF and Google.

* Corresponding author.

E-mail addresses: clifford@cs.bris.ac.uk (R. Clifford), klimefrem@gmail.com (K. Efremenko), bennyporat@gmail.com (B. Porat), porately@cs.biu.ac.il (E. Porat).

by recent advances for the problem of k -mismatch problem with do not cares [8,10] as well as a surge in interest in provably fast algorithms for distance calculation and approximate matching between numerical strings. Many different metrics have been considered, with, for example, the L_1 distance [4,5,7] and less-than matching [2] problems both being solvable in $O(n\sqrt{m \log m})$ time and a bounded version of the L_∞ norm which was first discussed in [9] and then improved in [7,21] requiring $O(\delta n \log m)$ time.

In almost every one of these cases and in many others beside, the algorithms make extensive use of the fast Fourier transform (FFT). The property of the FFT that is required is that in the RAM model, the cross-correlation,

$$(t \otimes p)[i] \stackrel{\text{def}}{=} \sum_{j=1}^m p_j t_{i+j-1}, \quad 0 \leq i \leq n - m + 1,$$

can be calculated accurately and efficiently in $O(n \log n)$ time (see e.g. [12, Chapter 32]). By a standard trick of splitting the text into overlapping substrings of length $2m$, the running time can be further reduced to $O(n \log m)$.

Although the FFT is a very powerful and successful tool, it also brings with it a number of disadvantages. Perhaps most significant of these in this context is that the cross-correlation computation using the FFT is very much an offline algorithm. It requires the entire pattern and text to be available before any search can be performed. Of course, it is not only the FFT that causes this difficulty. For example, the only fast algorithm for the k -mismatch algorithm which does not employ the FFT uses constant time LCA queries [20]. As it is not known how to perform the necessary preprocessing of a suffix tree to compute the LCA online, this algorithm suffers from the same limitations as those that depend on the FFT.

In many situations such as when monitoring Internet traffic or telecommunications networks this model of computation may not be feasible. It is not sufficient simply that a pattern matching algorithm runs fast. It should also require considerably less space than the input and update at least as quickly as the new data are arriving while still maintaining an overall time complexity which is as close as possible to the full offline algorithm. One approach to handle this situation is the data streaming model where it is assumed that it is not possible to ever store all the data seen and in some variants that only one pass over the data is ever allowed. This very successful model has been the source of a great deal of attention in recent years (see [15] and [22] for background on data stream computation), however the techniques developed have largely been randomised whereas our interest is in deterministic solutions.

We have also become aware at the final stage of preparation of this paper of work from 1973 on fast online integer multiplication [14]. The techniques contained therein have a considerable overlap with our basic approach of recursively splitting the pattern and also achieve the same log factor multiplicative overhead. However, this much earlier work is written in the different context of multitape Turing machines, focuses mostly on the problem of multiplication and does not consider the case where one input is much longer than the other as it is not concerned with pattern matching.

Our main contribution is a black box for converting offline approximate matching algorithms into efficient online ones. That is, it ensures that the approximate matching algorithm accomplishes its task for the i th input character without requiring the $i + 1$ th. The method is deterministic and bounds the worst case running time *per input character* as well as ensuring that the overall running time is within a log factor of the best known offline algorithm. It is an important feature of our method that its running time is not amortised. This is because when processing streaming data it may not be realistic to wait for long periods of time between individual input characters.

Our technique can be applied to a wide class of approximate matching algorithms overcoming one of the main restrictions on their use in data streaming applications. A particularly useful subset that we focus on in this paper includes problems for which the distance $\Delta(x, y)$ between two strings of length m can be defined as a simple function of $\sum_{j=1}^m \Delta(x_j, y_j)$. In other words, distance functions between two strings where the distance is measured using the sum of the distances between individual symbols. Many of the most common and widely studied approximate matching problems fall into this category including exact matching with wildcards, matching under the Hamming norm, k -mismatch and matching under the L_2 and L_1 norms. As a result, we provide fast deterministic online algorithms for each one of these problems.

The overall structure of the paper is as follows. In Section 2, we summarise the main results of the paper. In Section 3, we present the main black box solution and in Section 4 we discuss space lower bounds. Finally we conclude with some open problems in Section 5.

2. Our results

Our black box approach converts an offline approximate matching algorithm into efficient online algorithm. Let $T(n, m)$ be the total running time of the best known offline approximate matching algorithm for the problem being considered. The main results we present are as follows:

- We show how offline approximate matching algorithms can be turned into online algorithms with strict bounds on the computation time per input character. The main idea is to split the pattern into $O(\log m)$ subpatterns of successively halving length and to perform searches in parallel on carefully chosen partitions of the text. The partitions are chosen so that the work needed to compute the distance to a sliding window of the text is started $O(m)$ characters before it is needed.

The online algorithm takes $O(\sum_{j=1}^{\log_2 m} T(n, 2^{j-1})/n)$ time per input character. As $\sum_{j=1}^{\log_2 m} T(n, 2^{j-1}) \leq T(n, m) \log(m)$, this gives a near optimal deterministic solution to a wide class of online approximate matching problems. Where the time complexity of the offline algorithm is such that $T(n, m) \in \Omega(nm^\alpha)$ with $\alpha > 0$, as in, for example, pattern matching under the Hamming norm, then it follows from a standard geometric sum argument that there is no asymptotic time penalty at all caused by making the algorithm online.

- For data that arrives in an intermittent or bursty fashion, a small adjustment to the algorithm allows us to report the distance to a sliding window in the text in constant time after a new character arrives. Although the computation time per character is unchanged, we are able to move the majority of the work for future symbols until after a new symbol has been processed. This provides a solution for online approximate matching in a model where instant answers are needed once new data arrives.
- Applications of our black box method to exact matching with wildcards, matching under the Hamming norm, approximating the Hamming norm, k -mismatch and matching under the L_2 and L_1 norms result in algorithms that take $O(\log^2 m)$, $O(\sqrt{m \log m})$, $O(\log^2 m/\epsilon^2)$, $O(\sqrt{k \log k \log m})$, $O(\log^2 m)$ and $O(\sqrt{m \log m})$ time per character, respectively.
- Finally we argue that the space requirements for the online approximate matching problem are optimal in the deterministic setting for a large class of approximate pattern matching problems.

3. The black box for online approximate matching

The black box we present will make repeated calls to an offline approximate pattern matching algorithm which we call `offline-pm`. In order to simplify some of the explanation, we assume that the running time $T(n, m)$ of `offline-pm` can be expressed as $nT'(m)$ with $T'(m) \in O(m)$. This assumption is reasonable as the types of pattern matching problem we consider can all be solved naively in $O(nm)$ time. As a result of this simplification we have $O(n/mT(cm, m)) = O(T(n, m))$ for constant $c > 1$. We will also at times refer to a call to `offline-pm` as a search for the sake of brevity.

The basic idea of our black box is to split the pattern into $O(\log m)$ consecutive subpatterns each having half the length of the previous one. In this way $P_1 = p[1, \dots, m/2]$ and subpattern P_j has length $m2^{-j}$ for $1 \leq j \leq \log_2(m)$. $P_{\log_2(m)+1}$ is set to be the last character of the pattern. We then run `offline-pm` for each subpattern against the whole of the text. The distances found can then be added to an auxiliary array C . Specifically, for any subpattern starting at position j of the pattern, its distance to a substring starting at position i of the text will be added to the count at $C[i - j + 1]$. At the end of this step C will contain $\Delta(p, t[i, \dots, i + m - 1])$ for every location i in t . An example of the whole pattern aligned with the text and the contributions from each of the subpatterns to the Hamming distance is shown in Fig. 1.

This algorithm will call `offline-pm` $O(\log m)$ times and requires $O(n)$ extra space for the auxiliary array. The space requirement can be reduced to $O(m)$ by partitioning the text. For any subpattern of length m' , we partition the text into $n/(m' - 1)$ overlapping substrings of length $2m'$, each with an overlap of length m' with the previous partition. If we run `offline-pm` on each partition separately, the total time complexity over the whole text for each subpattern is $O((n/m)T(2m, m)) = O(T(n, m))$. The distances for each subpattern can be added to the auxiliary C in the same way as before. However, now we only need store one auxiliary array of size m at most.

This space reduced algorithm can be made online by a lazy execution of the searches performed. For each subpattern P_j there is a list of locations associated with it which marks out the start and end of its associated partitions of the text. These locations do not have to be stored explicitly as they are computable as they are needed. For example, for $P_1 = p[1, \dots, m/2]$, the first partition is $t[1, \dots, m]$, the second is $t[m/2, 3m/2]$ and so on. When the i th character is read in from the text, `offline-pm` can now perform the searches for all the subpatterns which have a partition finishing at position i . This online algorithm runs in $O(\sum_{j=1}^{\log_2 m} T(n, 2^{j-1}))$ time and $O(m)$ extra space. However the solution is not yet satisfactory as we might potentially have to wait for more than $T(m, m/2)$ time after a new character arrives before we are able to compute the distance to the new sliding window. As an example, when the m th character of t is read in, the first search involving P_1 commences for the partition $t[1, \dots, m]$, thereby delaying the computation of $\Delta(p, t[1, \dots, m])$ unacceptably. Our aim is to ensure that the maximum computation time per character is limited to $O(\sum_{j=1}^{\log_2 m} T(n, 2^{j-1})/n)$.

3.1. Bounding the maximum time per character

We can bound the maximum time per character by performing more work earlier on. Instead of splitting the text into partitions of size $2m'$ per subpattern, we change the partition size to $3m'/2$. The overlap between partitions is maintained at m' to ensure no matches are missed.




Fig. 1. Pattern splitting example using Hamming distance.




Fig. 2. Partitioning of the text for subpattern $P_j = abba$.

If the i th character of the text is read in, searches will now be performed for all subpatterns which have a partition of the text ending at position i . The result of a search involving a subpattern will now not be needed until the $i + m'/2$ th character is read in. That is $m'/2$ characters after the relevant search has been performed. In this way, whenever a new character is read in, the results for all the subpatterns needed to compute the distance to the new sliding window are already known, except for the last character of the pattern. This last comparison can be carried out in constant time. Fig. 2 shows the partitioning of the text for a subpattern P_j . In this example, the `offline-pm` algorithm will be called for this alignment of the subpattern P_j as soon as the character c at the end of the marked section of the text of length $3|P_j|/2$ is read in and the result will be required when another character c at the end the next section of length $|P_j|/2$ is inputted.

In order to guarantee the desired upper bound for the amount of work carried out per character, we need to spread the work out evenly over the time it takes to input the text. For each subpattern, the work for a particular search does not have to be completed until $m'/2$ characters after it starts and so we can set this work to be performed over the period between reading in the i th and the $i + m'/2$ th character. As an example, the first search for P_1 starts when $t[3m/4]$ is read in and completes at the time that $t[m]$ is. P_2 will on the other hand start to calculate its contribution to $\Delta(p, t[1, \dots, m])$ when $t[7m/8]$ is read in. The result from these searches and from all the other subpatterns, will be scheduled to complete by the time $t[m]$ arrives. As there will be a number of searches scheduled to work in this period we will need to perform the work in parallel by time slicing. To guarantee the bound on the work per character without requiring any internal knowledge of `offline-pm`, we only require an upper bound for the running time of `offline-pm` for a subpattern of length m' which we then divide by $m'/2$. This gives the amount of work to carry out per input character.

Algorithm 1 gives an overview of the whole process and the following Theorem summarises the main result. t

Input: Pattern p , the streaming text t and `offline-pm`
Output: $\Delta(p, t[i - m + 1, \dots, i])$ for each $i \geq m$ of streaming text
 Initialisation;
 Split p into $\log_2 m$ subpatterns P_j of length $m/2^j$, for $1 \leq j \leq \log_2 m$;
 For each P_j , calculate its partition start and end points;
foreach symbol $t[i]$ read in **do**
 Add $\Delta(p[m], t[i])$ to $C[i - m + 1]$;
 Wait for results of `offline-pm` searches due to end at position i ;
 Output $C[i - m + 1]$;
 Start `offline-pm` searches for each subpattern P_j which has a partition ending at i ;
end

Algorithm 1: Black box algorithm for online pattern matching.

Theorem 1. Algorithm 1 solves the online approximate problem in $O(\sum_{j=1}^{\log_2 m} T(n, 2^{j-1})/n)$ time per input character and $O(m)$ space.

Proof. The total time taken by Algorithm 1 is $O(T(n, m'))$ per subpattern of length m' making a total of $O(\sum_{j=1}^{\log_2 m} T(n, 2^{j-1}))$ time overall for all subpatterns. The work performed by the calls to `offline-pm` is evenly spread over the whole length of the text. Therefore, the total amount of work per character is $O(\sum_{j=1}^{\log_2 m} T(n, 2^{j-1})/n)$.

The space required for the auxiliary array is $O(m)$. We also have to consider the space overhead of `offline-pm` as there can be $O(\log m)$ searches running simultaneously. Under the assumption that each individual search requires $O(m)$ space, the total space requirement is bounded by $\sum_{j=1}^{\log_2 m} (c2^{j-1}) = O(m)$ overall. \square

The technique of scheduling work to start before it is needed and performed as new symbols are read in gives us a great deal of flexibility. For example, by making a small adjustment to the scheduling of the work, we can also guarantee that all but a constant amount of the work needed to compute $\Delta(p, t[i - m + 1, \dots, i])$ will have been completed *before* the i th character is read in. This trick is only applicable if there is enough time between the arrival of two characters to perform the

extra work needed. Although no change is made to the total work per character, the ability to control at which point work is carried out can have applications when the data arrives in bursts. For example, if there is some pause in the data stream before new characters arrive.

In order to output $\Delta(p, t[i - m + 1, \dots, i])$ in constant time after $t[i]$ is read in we need only ensure that a search is completed $m'/2 - 1$ (rather than $m'/2$) characters after the end of a partition. In this way, when the i th character is read in only $\Delta(p[m], t[i])$ will remain to be computed. This modification does not affect the time complexity of the algorithm overall.

We can now apply our black box to a number of well known matching problems, giving the following time complexities per input character.

Corollary 1. *Algorithm 1 applied to the fastest known offline pattern matching algorithms for the Hamming norm, k -mismatch and matching under the L_2 and L_1 norms gives online algorithms that take $O(\log^2 m)$, $O(\sqrt{m \log m})$, $O(\sqrt{k \log k \log m})$, $O(\log^2 m)$ and $O(\sqrt{m \log m})$ time per character, respectively.*

4. Space lower bound for deterministic online approximate matching

It would seem desirable to reduce the space requirements even further in order to increase the practicality of processing data streams. Unfortunately, we can show that this is not possible in the deterministic setting. Our argument follows from the simple observation that $m \log |\Sigma|$ bits are required in the worst case to send an arbitrary message of size m over an alphabet Σ between two parties, Alice and Bob. Under the assumption that the alphabet size $|\Sigma|$ is bounded above by the size of a word of memory, any online approximation pattern matching algorithm which requires $o(m)$ words of memory would allow us to send such a message using fewer bits, thereby giving a contradiction.

We assume only one property of the approximate pattern matching problem that is being considered. The property we require is that given a fixed pattern containing only standard symbols from the input alphabet (no wildcards, for example), the algorithm outputs a different answer when matched with another copy of the pattern than it does with any other string containing only standard symbols. To give the contradiction Alice starts the pattern matching algorithm by preprocessing the message which she wants to send to Bob. By our assumption the algorithm uses $o(m)$ words of space. As soon as the preprocessing is over, Alice transfers a snapshot of the current state of memory to Bob. Bob now continues the pattern matching process by trying every different string of standard symbols of length m in turn. When he finds an exact match he reports it and hence recovers the entirety of the message Alice wanted. The total number of words communicated is $o(m)$ thereby giving the desired contradiction.

5. Discussion

The method we have developed is applicable to a wide class of previously offline approximate matching algorithms. By choosing a black box approach we have not investigated whether particular pattern matching algorithms might be more directly converted to efficient online ones. For example, it is clear that the standard dynamic programming solution for edit distance can immediately be made online with no extra asymptotic time cost. However a bounded version of edit distance, the so called k -difference problem, where the fastest known offline solution uses the bound k in its running time, does not fit into our model and appears more challenging. Also, although we have shown that the space required by our approach is optimal for a wide range of problems, an interesting question is whether randomisation can allow us to solve the same problems with only $o(m)$ space as the communication complexity bounds will no longer hold.

Acknowledgment

The authors thank Inbok Lee and Ashley Montanaro for their helpful comments on a draft of this paper.

References

- [1] K. Abrahamson, Generalized string matching, *SIAM Journal on Computing* 16 (6) (1987) 1039–1051.
- [2] A. Amir, M. Farach, Efficient 2-dimensional approximate matching of half-rectangular figures, *Information and Computation* 118 (1) (1995) 1–11.
- [3] A. Amir, M. Lewenstein, E. Porat, Faster algorithms for string matching with k mismatches, *Journal of Algorithms* 50 (2) (2004) 257–275.
- [4] A. Amir, O. Lipsky, E. Porat, J. Uman, Approximate matching in the L_1 metric, in: *CPM'05: Proceedings of the 17th Annual Symposium on Combinatorial Pattern Matching*, Springer (2005) 91–103.
- [5] M.J. Atallah, Faster image template matching in the sum of the absolute value of differences measure, *IEEE Transactions on Image Processing* 10 (4) (2001) 659–663.
- [6] P. Clifford, R. Clifford, Simple deterministic wildcard matching, *Information Processing Letters* 101 (2) (2007) 53–54.
- [7] P. Clifford, R. Clifford, C.S. Iliopoulos, Faster algorithms for δ , γ -matching and related problems, in: *CPM'05: Proceedings of the 17th Annual Symposium on Combinatorial Pattern Matching*, Springer (2005) 68–78.
- [8] R. Clifford, K. Efremenko, E. Porat, A. Rothschild, k -Mismatch with don't cares, in: *ESA'07: Proceedings of the 15th Annual European Symposium on Algorithms* (October 2007) 151–162.
- [9] R. Clifford, C. Iliopoulos, String algorithms in music analysis, *Soft Computing* 8 (9) (2004) 597–603.

- [10] R. Clifford, E. Porat, A filtering algorithm for k -mismatch with don't cares, in: SPIRE'07: Proceedings of the 14th International Symposium on String Processing and Information Retrieval (October 2007) 130–136.
- [11] R. Cole, R. Hariharan, Verifying candidate matches in sparse and wildcard matching, in: STOC'02: Proceedings of the 34th Annual ACM Symposium Theory of Computing (2002) 592–601.
- [12] T.H. Cormen, C.E. Leiserson, R.L. Rivest, Introduction to Algorithms, MIT Press, 1990.
- [13] M. Fischer, M. Paterson, String matching and other products, in: R. Karp (Ed.), Proceedings of the 7th SIAM-AMS Complexity of Computation, 1971, pp. 113–125.
- [14] Michael J. Fischer, Larry J. Stockmeyer, Fast on-line integer multiplication, in: STOC'73: Proceedings of the 5th Annual ACM Symposium on Theory of Computing, 1973, pp. 67–72.
- [15] M.R. Henzinger, P. Raghavan, S. Rajagopalan, External Memory Algorithms, Chapter Computing on Data Streams, American Mathematical Society, Boston, MA, USA, 1999, pp. 107–118.
- [16] P. Indyk, Faster algorithms for string matching problems: matching the convolution bound, in: FOCS'98: Proceedings of the 39th Annual Symposium Foundations of Computer Science, 1998, pp. 166–173.
- [17] A. Kalai, Efficient pattern-matching with don't cares, in: SODA'02: Proceedings of the 13th ACM-SIAM Symposium on Discrete Algorithms, 2002, pp. 655–656.
- [18] H. Karloff, Fast algorithms for approximately counting mismatches, Information Processing Letters 48 (2) (1993) 53–60.
- [19] S.R. Kosaraju, Efficient string matching, Manuscript, 1987.
- [20] G.M. Landau, U. Vishkin, Efficient string matching with k mismatches, Theoretical Computer Science 43 (1986) 239–249.
- [21] O. Lipsky, E. Porat, Approximate matching in the L_∞ metric, in: SPIRE'05: Proceedings of the 12th International Symposium on String Processing and Information Retrieval, Springer (2005) 331–334.
- [22] S. Muthukrishnan, Data streams: algorithms and applications, Foundations and Trends in Theoretical Computer Science 1 (2) (2005).